

# Debugging and optimization of HPC programs with the Verrou tool

Software Correctness  
for HPC Applications  
18/11/2019


**François Févotte**  
**Bruno Lathuilière\***

EDF R&D  
PERICLES / I23  
Analysis and Numerical Modeling




# Industrial context – Numerical Verification


In-house development of Scientific Computing Codes




Structures




Fluid dynamics




Wave propagation



Neutronics



Power Systems



Free surface hydraulics

# Industrial context – Numerical Verification


## Code\_aster

### Mechanics

- ◆ Seismic
- ◆ Acoustic
- ◆ Thermo-mechanics

### Code\_Aster

- ◆ 1.5M code lines
- ◆ Fortran 90, C, Python
- ◆ thousands of test cases
- ◆ Large number of dependencies :
  - ▶ Linear solvers (MUMPS...)
  - ▶ Mesh generator and partitioning tools (Metis, Scotch...)



### Goals

- ◆ understand the non-reproducibility between test computers

# Industrial context – Numerical Verification

Objectives / presentation outline

## Diagnostics

- ◆ verify a code / show the presence of FP-related errors
- ◆ quantify the magnitude of issues

## Debugging

- ◆ locate the origin of FP-related issues in the source code
  - ▶ unstable algorithms
  - ▶ unstable tests
- ◆ track the origin of issues during program execution
  - ▶ context of calls
  - ▶ temporal information (e.g. iteration number...)

## Optimization

- ◆ use mixed-precision implementations



Available on github (latest version: v2.1.0)

<http://github.com/edf-hpc/verrou>

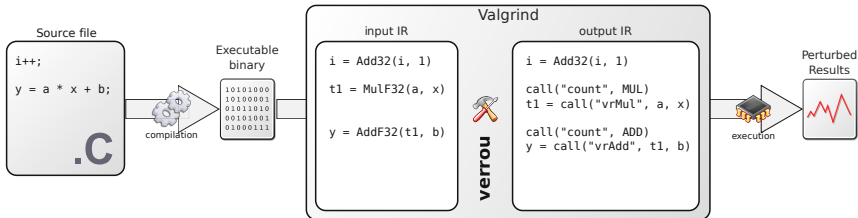
Documentation:

<http://edf-hpc.github.io/verrou/vr-manual.html>

# Industrial context – Numerical Verification

## Dynamic binary analysis with Valgrind


\$ valgrind --tool=verrou [VERROU\_ARGS] PROGRAM [ARGS...]



# Diagnostics: detect and assess instabilities

Verrou back-end: random rounding

IEEE-754: nearest rounding mode




# Diagnostics: detect and assess instabilities

Verrou back-end: random rounding

CESTAC: random rounding mode

`--rounding-mode=random`





# Basic use

## Run the code:

```
$ PROG='python -c "print(1+10**6-sum([0.1 for i in range(10**7)]))"'
$ eval $PROG
1.00016102463
```

## Verify the instrumentation process:

```
$ CMD="valgrind --tool=verrou --rounding-mode=nearest $PROG"
$ eval "$CMD" 2>/dev/null
1.00016102463
```

## Run stochastic samples:

```
$ CMD="valgrind --tool=verrou --rounding-mode=random $PROG"
$ eval "$CMD; $CMD; $CMD" 2>/dev/null
0.999830178451
0.999830617453
0.999830234447
```

## Post-process:

$$\hat{s} = -\log_2 \left( \frac{\max |X_i - X_{\text{ieee}}|}{|X_{\text{ieee}}|} \right)$$

$\approx 11.6$  significant bits  
(3.5 significant digits)

# Really few false positive

Libm is the noticeable exception

## Issue

```
$ PROG='python -c "import math;print(math.cos(42.))"'
$ CMD="valgrind --tool=verrou --rounding-mode=random $PROG"
$ eval "$CMD; $CMD; $CMD" 2>/dev/null
-0.399985314988
-0.399985314988
-1.00505077023
```

## Baseline: ignore FP operation inside libm

```
$ cat libm.ex
* /lib/x86_64-linux-gnu/libm-2.19.so
$ CMD="valgrind --tool=verrou --rounding-mode=random --exclude=libm.ex $PROG"
$ eval "$CMD; $CMD; $CMD" 2>/dev/null
-0.399985314988
-0.399985314988
-0.399985314988
```

# Interlibm

An interposition library

Use:

```
VERROU_ROUNDING_MODE=random LD_PRELOAD=$VERROU_LIBM $PROG
```

Implementation for each libm function:

- ◆ To reuse the verrou rounding mode simulation we need:
  - ▶ the round to nearest result
  - ▶ and an estimation of error

The implementation uses the estimation provided by libquadmath.

## Exemple extracted from code `_aster`

for  $a = 4.208003496301644 \times 10^{-5}$  and  $b = a + 6 \text{ ulp}(a)$ .

$$f(a, b) = \frac{b - a}{\log(b) - \log(a)}$$

Number of significant bits for both implementations.

	$\frac{b - a}{\log(b) - \log(a)}$	$a \frac{\frac{b}{a} - 1}{\log(\frac{b}{a})}$
IEEE Error	16.60	53.00
(i) interlibm	14.12	52.46
(ii) verrou	52.46	51.46
(iii) verrou+interlibm	14.12	50.88

Estimation with 459 samples and following estimator :

$$\hat{s} = -\log_2 \left( \frac{\max |X_i - X_{\text{ieee}}|}{|X_{\text{ieee}}|} \right). \quad (1)$$



# Plan

1. Diagnostics
2. Debugging
3. Conclusions – perspectives



# Debugging: code coverage

## Instable tests detection

```
$ make CFLAGS="-fprofile-arcs -ftest-coverage"  
$ make check  
$ gcov *.c *.f
```

### "standard" coverage

```
120:subroutine fun1(area, a1, a2, n)  
-:   implicit none  
-:   integer :: n  
-:   real(kind=8) :: area, a1, a2  
120:   if (a1 .eq. a2) then  
13:       area = a1  
-:   else  
107:       if (n .lt. 2) then  
107:           area = (a2-a1) / (log(a2)-log(a1))  
###:       else if (n .eq.2) then  
###:           area = sqrt (a1*a2)  
-:       else  
###:           ! ...  
-:       endif  
-:   endif  
120:end subroutine
```

# Debugging: code coverage

## Instable tests detection

```
$ make CFLAGS="-fprofile-arcs -ftest-coverage"  
$ make check  
$ gcov *.c *.f
```

### "standard" coverage

```
120:subroutine fun1(area, a1, a2, n)  
-:    implicit none  
-:    integer :: n  
-:    real(kind=8) :: area, a1, a2  
120:    if (a1 .eq. a2) then  
13:        area = a1  
-:    else  
107:        if (n .lt. 2) then  
107:            area = (a2-a1) / (log(a2)-log(a1))  
###:        else if (n .eq.2) then  
###:            area = sqrt (a1*a2)  
-:        else  
###:            ! ...  
-:        endif  
-:    endif  
120:end subroutine
```

Debugging and optimization of HPC programs with the Verrou tool

### "Verrou" coverage

```
120:subroutine fun1(area, a1,...  
-:    implicit none  
-:    integer :: n  
-:    real(kind=8) :: area,...  
120:    if (a1 .eq. a2) then  
4:        area = a1  
-:    else  
116:        if (n .lt. 2) then  
116:            area = (a2-a1...  
###:        else if (n .eq.2)...  
###:            area = sqrt (...  
-:        else  
###:            ! ...  
-:        endif  
-:    endif  
120:end subroutine
```

# Debugging: locate issues in the source code

## Delta-debugging

log_L	.../aster.release
volum2_	.../aster.release
bilpla_	.../aster.release
ecrval_	.../aster.release
print_plath_	.../aster.release
classer_groupes_	.../aster.release
etupla_	.../aster.release
couhyd_pi_	.../aster.release
ecrplr_	.../aster.release
imovi_	.../aster.release
resopt_	.../aster.release
getgrp_marginal_	.../aster.release
ecrpla_	.../aster.release
fin_exec_main_	.../aster.release
decopt_pi_	.../aster.release
paraend_	.../aster.release
resopt_cnt_zones_	.../aster.release
apstop_	.../aster.release
ihyd_	.../aster.release
impression_info_	.../aster.release
coupla_	.../aster.release
gere_print_plath_	.../aster.release
log	.../aster.release
thepla_	.../aster.release
coutot_	.../aster.release
iprit_	.../aster.release

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations





# Debugging: locate issues in the source code

## Delta-debugging

```
# log.L                .../aster.release
# volum2_              .../aster.release
# bilpla_              .../aster.release
# ecrval_              .../aster.release
# print_plath_         .../aster.release
# classer_groupes_     .../aster.release
# etupla_              .../aster.release
# couhyd_pi_           .../aster.release
# ecrplr_              .../aster.release
# imovi_               .../aster.release
# resopt_              .../aster.release
# getgrp_marginal_     .../aster.release
# ecrpla_              .../aster.release
# fin_exec_main_       .../aster.release
# decopt_pi_           .../aster.release
# paraend_             .../aster.release
# resopt_cnt_zones_    .../aster.release
# apstop_              .../aster.release
# ihyd_                .../aster.release
# impression_info_     .../aster.release
# coupla_              .../aster.release
# gere_print_plath_    .../aster.release
# log                  .../aster.release
# thepla_              .../aster.release
# coutot_              .../aster.release
# iprit_               .../aster.release
```

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations



# Debugging: locate issues in the source code

## Delta-debugging

```
# log.L          .../aster.release
# volum2_       .../aster.release
# bilpla_       .../aster.release
# ecrval_       .../aster.release
# print_plath_  .../aster.release
# classer_groupes_ .../aster.release
# etupla_       .../aster.release
# couhyd_pi_    .../aster.release
# ecrplr_       .../aster.release
# imovi_        .../aster.release
# resopt_       .../aster.release
# getgrp_marginal_ .../aster.release
# ecrpla_       .../aster.release
fin_exec_main_  .../aster.release
decopt_pi_     .../aster.release
paraend_       .../aster.release
resopt_cnt_zones_ .../aster.release
apstop_        .../aster.release
ihyd_          .../aster.release
impression_info_ .../aster.release
coupla_        .../aster.release
gere_print_plath_ .../aster.release
log            .../aster.release
thepa_         .../aster.release
coutot_        .../aster.release
iprit_         .../aster.release
```

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations
- ◆ Delta-Debugging [A. Zeller, 1999] adapted for stochastic evaluation



# Debugging: locate issues in the source code

## Delta-debugging

```
# log.L          .../aster.release
# volum2_       .../aster.release
# bilpla_       .../aster.release
# ecrval_       .../aster.release
# print_plath_  .../aster.release
# classer_groupes_ .../aster.release
# etupla_       .../aster.release
couhyd_pi_      .../aster.release
ecrplr_         .../aster.release
imovi_         .../aster.release
resopt_        .../aster.release
getgrp_marginal_ .../aster.release
ecrpla_        .../aster.release
fin_exec_main_  .../aster.release
decopt_pi_     .../aster.release
paraend_       .../aster.release
resopt_cnt_zones_ .../aster.release
apstop_        .../aster.release
ihyd_          .../aster.release
impression_info_ .../aster.release
coupla_        .../aster.release
gere_print_plath_ .../aster.release
log            .../aster.release
thepa_         .../aster.release
coutot_        .../aster.release
iprit_         .../aster.release
```

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations
- ◆ Delta-Debugging [A. Zeller, 1999] adapted for stochastic evaluation



# Debugging: locate issues in the source code

## Delta-debugging

```
log.L                .../aster.release
volum2_              .../aster.release
bilpla_              .../aster.release
ecrval_              .../aster.release
print_plath_         .../aster.release
classer_groupes_     .../aster.release
etupla_              .../aster.release
# couhyd_pi_         .../aster.release
# ecrplr_             .../aster.release
# imovi_             .../aster.release
# resopt_            .../aster.release
# getgrp_marginal_   .../aster.release
# ecrpla_            .../aster.release
fin_exec_main_       .../aster.release
decopt_pi_           .../aster.release
paraend_             .../aster.release
resopt_cnt_zones_    .../aster.release
apstop_              .../aster.release
ihyd_                .../aster.release
impression_info_     .../aster.release
coupla_              .../aster.release
gere_print_plath_    .../aster.release
log                  .../aster.release
thepla_              .../aster.release
coutot_              .../aster.release
iprit_               .../aster.release
```

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations
- ◆ Delta-Debugging [A. Zeller, 1999] adapted for stochastic evaluation



# Debugging: locate issues in the source code

## Delta-debugging

```
log.L                .../aster.release
volum2_              .../aster.release
bilpla_              .../aster.release
ecrval_              .../aster.release
print_plath_         .../aster.release
classer_groupes_     .../aster.release
etupla_              .../aster.release
# couhyd_pi_         .../aster.release
ecrplr_              .../aster.release
imovi_               .../aster.release
resopt_              .../aster.release
getgrp_marginal_     .../aster.release
ecrpla_              .../aster.release
fin_exec_main_       .../aster.release
# decopt_pi_         .../aster.release
paraend_             .../aster.release
resopt_cnt_zones_    .../aster.release
apstop_              .../aster.release
# ihyd_              .../aster.release
impression_info_     .../aster.release
coupla_              .../aster.release
gere_print_plath_    .../aster.release
log                  .../aster.release
thepla_              .../aster.release
# coutot_            .../aster.release
# iprit_              .../aster.release
```

- ◆ relies on the Verrou ability to restrict the scope of instrumentation/perturbations
- ◆ Delta-Debugging [A. Zeller, 1999] adapted for stochastic evaluation
- ◆ Inputs :
  - ▶ run script
  - ▶ comparison script
- ◆ Output:
  - ▶ “unstable” code parts
- ◆ Also works at the source line granularity:
  - ▶ if the code was compiled with `-g`

# Conclusions

Verrou as a tool to help with FP issues

## Diagnostics

- ☀ show the presence of FP-related errors (random-rounding back-end)
- 🌱 quantify the magnitude of issues (post-processing)

## Debugging

- 🌱 locate the origin of FP-related issues in the source code
  - ☀ unstable algorithms (Delta-Debugging)
  - ☁ unstable tests (code coverage analysis)
- ⚡ track the origin of issues during program execution
  - ▶ context of calls
  - ▶ temporal information (e.g. iteration number...)

## Optimization

- ☁ emulate mixed-precision implementations (reduced precision back-end)
- 🔱 re-use debugging features (Delta-Debugging)

# Outlooks

## Interflop

- ◆ Verrou is no silver bullet
  - ▶ multiply techniques & tools

## Interflop (toolbox)

Common interface for Verificarlo & Verrou

- ◆ share Stochastic Arithmetic back-ends
- ◆ share accompanying tools (Delta-Debugging...)
- ◆ improve performance of instrumentation front-ends

## Interflop (larger consortium)

- ◆ Explore different analysis methods and the links between them
  - ▶ Stochastic Arithmetic
  - ▶ Interval Arithmetic & Affine Forms



Thank you !  
Questions ?

Get Verrou on github:  
<http://github.com/edf-hpc/verrou>

Documentation:  
<http://edf-hpc.github.io/verrou/vr-manual.html>



# Relevant references I



Jean-Marie Chesneaux and Jean Vignes, *On the robustness of the cestac method*, C. R. Acad.Sci. Paris **1** (1988), 855–860.



Christophe Denis, Pablo de Oliveira Castro, and Eric Petit, *Verificarlo: checking floating point accuracy through Monte Carlo Arithmetic*, 23rd IEEE Internatinal Symposium on Computer Arithmetic (ARITH'23), 2016.



François Févotte and Bruno Lathuilière, *VERROU: Assessing Floating-Point Accuracy Without Recompiling*,  
<https://hal.archives-ouvertes.fr/hal-01383417>, October 2016.



François Févotte and Bruno Lathuilière, *Studying the numerical quality of an industrial computing code: A case study on code\_aster*, 10th International Workshop on Numerical Software Verification (NSV) (Heidelberg, Germany), July 2017, pp. 61–80.

## Relevant references II



Stef Graillat, Fabienne Jézéquel, and Romain Picot, *Numerical validation of compensated algorithms with stochastic arithmetic*, Applied Mathematics and Computation **329** (2018), 339 – 363.



Fabienne Jézéquel, Jean-Marie Chesneaux, and Jean-Luc Lamotte, *A new version of the CADNA library for estimating round-off error propagation in Fortran programs*, Computer Physics Communications **181** (2010), no. 11, 1927–1928.







William Kahan, *How futile are mindless assessments of roundoff in floating-point computations?*, <https://people.eecs.berkeley.edu/~wkahan/Mindless.pdf>, 2006.



Jean-Luc Lamotte, Jean-Marie Chesneaux, and Fabienne Jézéquel, *CADNA\_C: A version of CADNA for use with C or C++ programs*, Computer Physics Communications **181** (2010), no. 11, 1925–1926.

## Relevant references III

-  Devan Sohier, Pablo De Oliveira Castro, François Févotte, Bruno Lathuilière, Eric Petit, and Olivier Jamond, *Confidence Intervals for Stochastic Arithmetic*, preprint, <https://hal.archives-ouvertes.fr/hal-01827319>.
-  Douglas Stott Parker, *Monte Carlo arithmetic: exploiting randomness in floating-point arithmetic*, Tech. Report CSD-970002, University of California, Los Angeles, 1997.
-  Jean Vignes, *A stochastic arithmetic for reliable scientific computation*, *Mathematics and Computers in Simulation* **35** (1993), 233–261.
-  Andreas Zeller, *Yesterday, My Program Worked. Today, It Does Not. Why?*, *SIGSOFT Softw. Eng. Notes* **24** (1999), no. 6, 253–267.



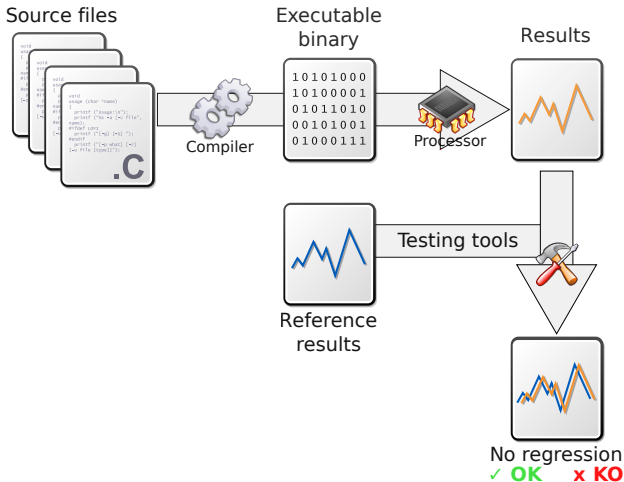
## Annexes

- ① Verrou
- ② Application : Athena

# The Verrou tool

## Development and QA process

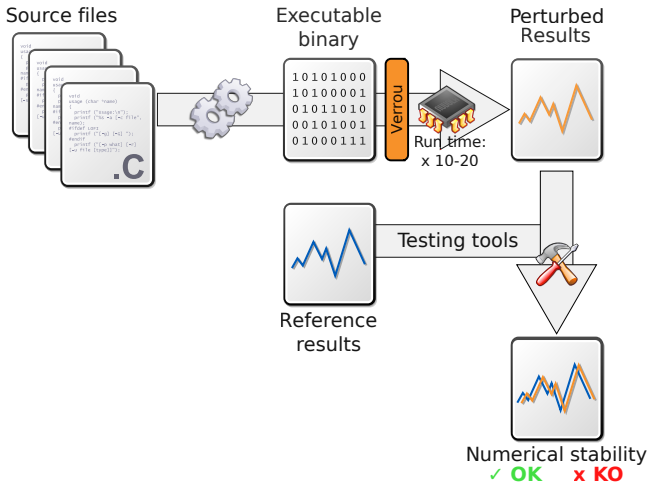
\$ myProg in out



# The Verrou tool

The Verrou tool: test the robustness w.r.t changes in the arithmetic

\$ `valgrind --tool=verrou --rounding-mode=random myProg in out`



# The Verrou tool

## Output exemple



```
$ valgrind --tool=verrou --rounding-mode=random PROGRAM [ARGS...]
```

```
==4683== Verrou, Check floating-point rounding errors
==4683== Copyright (C) 2014, F. Fevotte & B. Lathuiliere.
...
==4683== First seed : 1430818339
==4683== Simulating AVERAGE rounding mode
==4683== Instrumented operations :
==4683==   add : yes
...
==4683== -----
==4683== Operation                      Instructions count
==4683==   '- Precision      Instrumented          Total
==4683== -----
==4683== add                  500869335              500869335      (100%)
==4683==   '- flt              400695468              400695468      (100%)
==4683==   '- dbl              100173867              100173867      (100%)
==4683== -----
==4683== sub                  763127658              763127658      (100%)
==4683==   '- flt              763127658              763127658      (100%)
==4683== -----
==4683== mul                  1202086563              1202086563      (100%)
==4683==   '- flt              1101912537              1101912537      (100%)
==4683==   '- dbl              100174026              100174026      (100%)
==4683== -----
...
```

# The Verrou tool

## Output exemple

```
$ valgrind --tool=verrou --rounding-mode=random PROGRAM [ARGS...]
```

```
==4683== Verrou, Check floating-point rounding errors
==4683== Copyright (C) 2014, F. Fevotte & B. Lathuilliere.
```

```
...
```

```
==4683== First seed : 1430818339
==4683== Simulating AVERAGE rounding mode
==4683== Instrumented operations :
```

```
==4683== add : yes
```

```
...
```

```
==4683== normal output of the program
```

```
==4683==
==4683== + Warnings for “dangerous” instructions
```

(ex : x87)

```
==4683==                                     (100%)
==4683==                                     (100%)
==4683==                                     (100%)
==4683== - dbl 100173867 100173867
```

```
==4683== -----
==4683== sub 763127658 763127658 (100%)
==4683== '- flt 763127658 763127658 (100%)
==4683== -----
```


```
==4683== mul 1202086563 1202086563 (100%)
==4683== '- flt 1101912537 1101912537 (100%)
==4683== '- dbl 100174026 100174026 (100%)
==4683== -----
```

```
...
```



# The Verrou tool

## Rounding mode simulation



- ◆ Error Free Transformation  
(the division is more complicated):


- ▶  $a \circ b = \sigma + \delta$ ,
- ▶  $\sigma = fl(a \circ b)$

- ◆ If  $\delta < 0$  :

- ▶  $\lfloor a \circ b \rfloor = fl(a \circ b) - ulp$ ,
- ▶  $\lceil a \circ b \rceil = fl(a \circ b)$ .

# The Verrou tool

## Rounding mode simulation



### ◆ Error Free Transformation

(the division is more complicated):


- ▶  $a \circ b = \sigma + \delta$ ,
- ▶  $\sigma = fl(a \circ b)$

### ◆ If $\delta > 0$ :

- ▶  $\lfloor a \circ b \rfloor = fl(a \circ b)$ ,
- ▶  $\lceil a \circ b \rceil = fl(a \circ b) + ulp$ .

# The Verrou tool

## Rounding mode simulation




### random (CESTAC)

- ▶  $p(\lceil a \circ b \rceil) = 0.5$
- ▶  $p(\lfloor a \circ b \rfloor) = 0.5$



# The Verrou tool

## Rounding mode simulation



### ► average

$$\text{► } p(\lceil a \circ b \rceil) = \frac{\delta}{ulp} \quad (\text{if } \delta > 0)$$

$$\text{► } p(\lfloor a \circ b \rfloor) = \frac{ulp - \delta}{ulp}$$



# Approximate transformation for division



◆ What we want:

$$\frac{a}{b} = q + r,$$

with  $q = \text{fl}(a/b)$ .

◆ proposed algorithm:

**Input** :  $a, b$

**Output** :  $\tilde{r}$  such as

$$a/b \simeq \text{fl}(a/b) + \tilde{r}.$$

- 1  $q \leftarrow \text{fl}(a/b)$
- 2  $(p, s) \leftarrow \text{twoprod}(b, q)$
- 3  $t \leftarrow \text{fl}(a - p)$
- 4  $u \leftarrow \text{fl}(t - s)$
- 5  $\tilde{r} \leftarrow \text{fl}(u/b)$

◆ Idea of proof :

$$q = \frac{a}{b} (1 + \epsilon_1)$$

$$\begin{aligned} p &= b q (1 + \epsilon_2) \\ &= a (1 + \epsilon_1) (1 + \epsilon_2) \end{aligned}$$

$$t = a - p \quad (\text{Sterbenz lemma})$$

$$\begin{aligned} u &= (t - s) (1 + \epsilon_3) \\ &= \left( a - (p + s) \right) (1 + \epsilon_3) \\ &= (a - bq) (1 + \epsilon_3) \\ &= b r (1 + \epsilon_3) \end{aligned}$$

$$\begin{aligned} \tilde{r} &= \frac{u}{b} (1 + \epsilon_4) \\ &= r (1 + \epsilon_3) (1 + \epsilon_4). \end{aligned}$$

# The Verrou tool

## Using Verrou and Random Rounding



Test case	nearest
ssls108i	OK
ssls108j	OK
ssls108k	OK
ssls108l	OK
sdnl112a	OK
ssnp130a	OK
ssnp130b	OK
ssnp130c	OK
ssnp130d	OK

# The Verrou tool

## Using Verrou and Random Rounding



Test case	nearest	rnd <sub>1</sub>
ssls108i	OK	OK
ssls108j	OK	OK
ssls108k	OK	OK
ssls108l	OK	OK
sdnl112a	OK	KO
ssnp130a	OK	OK
ssnp130b	OK	OK
ssnp130c	OK	OK
ssnp130d	OK	OK

# The Verrou tool

## Using Verrou and Random Rounding



Test case	Status			
	nearest	rnd <sub>1</sub>	rnd <sub>2</sub>	rnd <sub>3</sub>
ssls108i	OK	OK	OK	OK
ssls108j	OK	OK	OK	OK
ssls108k	OK	OK	OK	OK
ssls108l	OK	OK	OK	OK
sdnl112a	OK	KO	KO	KO
ssnp130a	OK	OK	OK	OK
ssnp130b	OK	OK	OK	OK
ssnp130c	OK	OK	OK	OK
ssnp130d	OK	OK	OK	OK

10 minutes

20 minutes each

(72 test cases)



# The Verrou tool

## Using Verrou and Random Rounding



Test case	nearest	Status rnd <sub>1</sub>	rnd <sub>2</sub>	rnd <sub>3</sub>	# common decimal digits C(rnd <sub>1</sub> , rnd <sub>2</sub> , rnd <sub>3</sub> )
ssls108i	OK	OK	OK	OK	10
ssls108j	OK	OK	OK	OK	10
ssls108k	OK	OK	OK	OK	10
ssls108l	OK	OK	OK	OK	9
sdnl112a	OK	KO	KO	KO	3
ssnp130a	OK	OK	OK	OK	9
ssnp130b	OK	OK	OK	OK	9
ssnp130c	OK	OK	OK	OK	9
ssnp130d	OK	OK	OK	OK	9

10 minutes

(72 test cases)

20 minutes each

$$C(x) = \log_{10} \left| \frac{\mu(x)}{\sigma(x)} \right|$$

# The Verrou tool

## Delta-Debugging variants and improvements



### Definition of “unstable” program parts

- ▶ DDmax: straws that breaks the camels back
- ▶ DDmin: parts that cause instabilities on their own

### Preliminary filtering of the list

- ▶ only consider functions performing FP computations
  - ▶ code\_aster test case (tt1v300a): 4459 → 154 symbols (96.5% reduction)
- ▶ a preliminary basic search might be beneficial for such small lists

### Account for the stochastic nature of tests

- ▶ progressively increase the number of correct runs required to validate a set

# The Verrou tool

## Comparison of Delta-Debugging algorithms



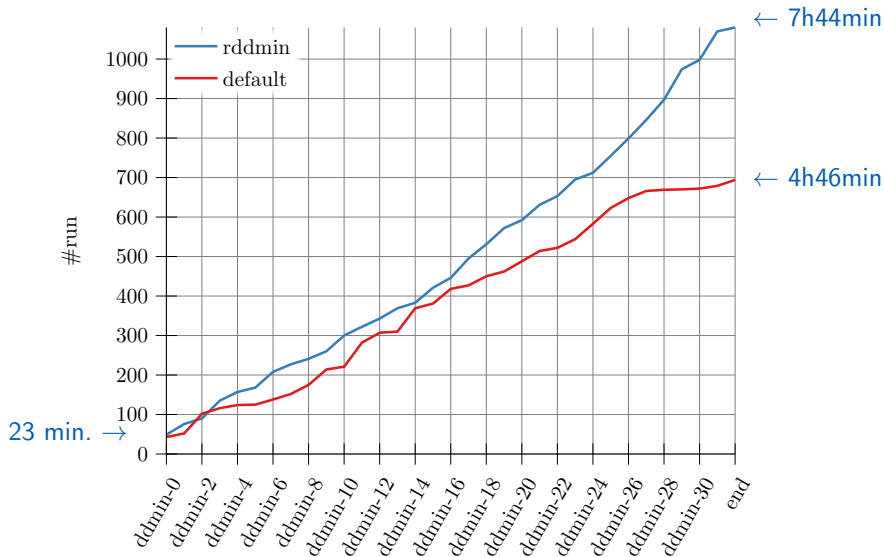
- ◆ Code\_aster test case (ttl300a)
  - ▶ 10 correct runs to validate a set
  - ▶ pre-filtering of FP functions

Algo	Total time	Speed-up
DDmax	12h 57min	
rDDmin	7h 44min	1.7
new default	4h 46min	2.7

- ◆ Tested algorithms
  - ▶ standard DDmax
  - ▶ standard DDmin
  - ▶ Verrou's improved DDmin (new default)
    - ▶ preliminary binary search with 5 correct runs requirement
    - ▶ exponential progression of the number of runs (1,2,5,10)



# The Verrou tool

## Standard DDmin algorithm vs Verrou's new default




# Application : NDT by US

Code Athena 2D




Produced result : "A-scan"



# Application : NDT by US

Code Athena 2D




Non destructive testing by Ultra-Sonic technics

## Code Athena 2D

- ◆ 36k code lines
- ◆ Fortran 77 + Fortran 90
- ◆ Dependencies : BLAS, LAPACK
- ◆ Code “known”

## Goal

- ◆ No identified problem
- ◆ “routine” test



# Application : NDT by US

Non-regression test with verrou



random 1	
<b>Cas-Test A</b>	
ins1.dat	0
ascan.dat	1.8e-12
<b>Cas-Test B</b>	
sismo.dat	7.9e-69
ascan.dat	1.2e-10
<b>Cas-Test C</b>	
ins1.dat	4.6e-06
sismo.dat	8.0e-28
ascan.dat	2.0e-11
<b>Cas-Test D</b>	
ins1.dat	1.5e-18
enerloc.dat	0
sismo.dat	0
ascan.dat	0

# Application : NDT by US

Non-regression test with verrou



	random 1	random 2
<b>Cas-Test A</b>		
ins1.dat	0	6.1e-06
ascan.dat	1.8e-12	5.9e-12
<b>Cas-Test B</b>		
sismo.dat	7.9e-69	7.9e-69
ascan.dat	1.2e-10	2.0e-11
<b>Cas-Test C</b>		
ins1.dat	4.6e-06	4.6e-06
sismo.dat	8.0e-28	2.8e-28
ascan.dat	2.0e-11	1.2e-11
<b>Cas-Test D</b>		
ins1.dat	1.5e-18	4.1e-01
enerloc.dat	0	2.3e-01
sismo.dat	0	1.6e-01
ascan.dat	0	1.5e-01



# Application : NDT by US


Non-regression test with verrou



	random 1	random 2	random 3	random 4
<b>Cas-Test A</b>				
ins1.dat	0	6.1e-06	6.1e-06	6.1e-06
ascan.dat	1.8e-12	5.9e-12	5.9e-12	5.9e-12
<b>Cas-Test B</b>				
sismo.dat	7.9e-69	7.9e-69	4.3e-69	4.3e-69
ascan.dat	1.2e-10	2.0e-11	2.8e-10	1.1e-11
<b>Cas-Test C</b>				
ins1.dat	4.6e-06	4.6e-06	4.6e-06	0
sismo.dat	8.0e-28	2.8e-28	8.0e-28	0
ascan.dat	2.0e-11	1.2e-11	1.8e-11	0
<b>Cas-Test D</b>				
ins1.dat	1.5e-18	4.1e-01	2.0e-01	0
enerloc.dat	0	2.3e-01	1.2e-01	0
sismo.dat	0	1.6e-01	3.2e-02	0
ascan.dat	0	1.5e-01	3.6e-01	6.5e-03


# Application : NDT by US

Test case D



# Application : NDT by US

Test case D



# Application : NDT by US

## Verrou performances



	reference	random	average
Cas-Test A	4.70s	83.23s (x17)	90.49s (x19)
Cas-Test B	29.79s	969.54s (x32)	1042.02s (x34)
Cas-Test C	21.15s	326.81s (x15)	358.08s (x16)
Cas-Test D	1.99s	24.20s (x12)	25.87s (x12)
Cas-Test E	0.46s	7.88s (x17)	8.88s (x19)
Cas-Test F	0.38s	4.54s (x11)	4.95s (x12)
Cas-Test G	6.16s	100.31s (x16)	109.70s (x17)
Cas-Test H	14.09s	503.90s (x35)	549.50s (x39)
Cas-Test I	1.48s	14.34s (x9)	14.85s (x10)

Slow down between  $\times 9$  and  $\times 39$



### ◆ Performance is the primary concern


- ▶ if more **powerful hardware** can be used, it will be used, ...
- ▶ even if **new developments** are needed to harness such hardware
- ▶ and **no holds are barred**:
  - ▶ as low as possible precision
  - ▶ unsafe compiler optimizations (such as `-ffast-math`)
  - ▶ loop vectorization / unrolling / both
  - ▶ hand-tuned assembly code
  - ▶ ...

### ◆ Every high-performance developer knows these techniques are unsafe


- ▶ and probably learned the hard way,
- ▶ but it is part of the trade.

# Additional slides

## V&V process: ad-hoc numerical instability detection methods



Physical input: affects the result  
Simulation parameter: should be neutral



► Idea: measure the sensitivity of the results w.r.t “neutral” parameters




easy to do



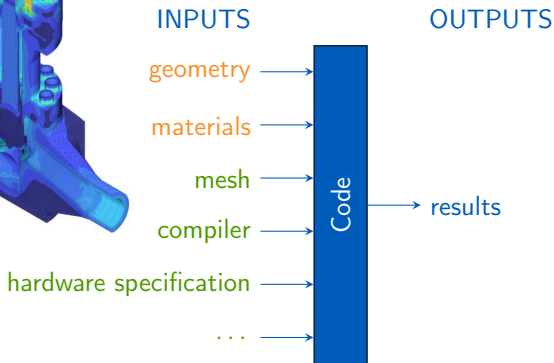
*ad hoc*, no localization

# Additional slides

## V&V process: ad-hoc numerical instability detection methods




Physical input: affects the result  
Simulation parameter: should be neutral



# Additional slides


## Example 1: failing computation in free surface hydraulics (Telemac 2D)






# Additional slides

## Example 1: failing computation in free surface hydraulics (Telemac 2D)



# Additional slides

## Example 1: failing computation in free surface hydraulics (Telemac 2D)



# Additional slides

## Example 2: performances of an optimization problem



### Initial problem formulation

$$\max_{p,v} \sum_{t \in T} \sum_{i \in I} \lambda_t p_t^i + \sum_{r \in R} \omega_r \left( v_t^r - V_0^r - \sum_{t \in T} \Gamma_t^r \right)$$

### Objective functional computation

$$\sum_{t \in T} \sum_{i \in I} \lambda_t p_t^i + \sum_{r \in R} \omega_r v_t^r \quad 1534019.677371745$$

$$- \sum_{r \in R} \omega_r \left( V_0^r + \sum_{t \in T} \Gamma_t^r \right) \quad -1534019.677282780$$

0.00008896500000000000

11 digits

# Additional slides

## Example 2: performances of an optimization problem



### Proposed re-formulation

$$\left[ \max_{p,v} \sum_{t \in T} \sum_{i \in I} \lambda_t p_t^i + \sum_{r \in R} \omega_r v_t^r \right] - \sum_{r \in R} \omega_r \left( V_0^r + \sum_{t \in T} \Gamma_t^r \right)$$

### Objective functional computation

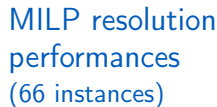
$$\sum_{t \in T} \sum_{i \in I} \lambda_t p_t^i + \sum_{r \in R} \omega_r v_t^r \quad 1534019.677371745$$

$$- \sum_{r \in R} \omega_r \left( V_0^r + \sum_{t \in T} \Gamma_t^r \right) \quad -1534019.677282780$$

---

0.00008896500000000000

11 digits



- 2 (3%) slowed down ( $\times 2$ )
- 15 (23%) unchanged
- 49 (74%) speed up ( $\times 2$ –1100)

# Additional slides



## Case study on ACTS: emulation of reduced precision + NaN detection

```
$ valgrind --tool=verrou --rounding-mode=float --exclude=excludes.ex --demangle=no --trace-children=yes --num-callers=50 \  
> ctest -V -R "BFieldMapUtils"
```

[...]

==1863== NaN:

```
==1863== at 0x5847E7F: _ZN5Eigen8internal4pmulIDv2_dEET_RKS3_S5_(emmintrin.h:271)  
==1863== by 0x585B570: _ZNK5Eigen8internal17scalar_product_opIddE8packetOpIDv2_dEET_RS6_S7_(BinaryFuncctors.h:89)  
==1863== by 0x59A31C9: _ZNK5Eigen8internal16binary_evaluatorINS_13CwiseBinaryOpINSO_17scalar_product_opIddE8KNS_14CwiseNullaryOpINSO_18scalar_constant_opIddE8KNS_6MatrixIdLi2ELi1ELi0ELi2ELi1EEEEESA_EENS0_10IndexBasedESE_ddE6packetILi16EDv2_dEETO_11(CoreEvaluators.h:727)
```

[...]

```
==1863== by 0x5BDE2D4: _ZN5Eigen6MatrixIdLi2ELi1ELi0ELi2ELi1EEaSINS_13CwiseBinaryOpINS_8internal13scalar_sum_opIddE8KNS3_INS4_17scalar_product_opIddE8KNS_14CwiseNullaryOpINS4_18scalar_constant_opIddE8KNS1_EESC_EESG_EEEERS1_RKNS_9DenseBaseIT_EE(Matrix.h:225)  
==1863== by 0x5BDAD82: _ZN4Acts6detail16interpolate_implIN5Eigen6MatrixIdLi2ELi1ELi0ELi2ELi1EEEE4_St5arrayIdLm2EES6_Lm1ELm4EE3runERKS4_RKS6_SB_RKS5_IS4_Lm4EE(interpolation_impl.hpp:133)  
==1863== by 0x5BD59E8: _ZN4Acts11interpolateIN5Eigen6MatrixIdLi2ELi1ELi0ELi2ELi1EEELm4ES3_St5arrayIdLm2EES5_vEET_RKT1_RKT2_RKT3_RKS4_IS6_XTO_EE(Interpolation.hpp:95)
```

[...]

```
==1863== by 0x4F6D54: _ZNK4Acts21InterpolatedBFieldMap11FieldMapperILj2ELj2EE8getFieldERKN5Eigen6MatrixIdLi3ELi1ELi0ELi3ELi1EEEE(InterpolatedBFieldMap.hpp:166)  
==1863== by 0x477030: _ZN4Acts4Test15bfield_creation11test_methodEv(BFieldMapUtilsTests.cpp:88)  
==1863== by 0x474767: _ZN4Acts4TestL23bfield_creation_invokerEv(BFieldMapUtilsTests.cpp:25)  
==1863== by 0x549B9B: _ZN5boost6detail18function22void_function_invokerOIPFvvEvE6invokeERNs1_15function_bufferE(function_template.hpp:118)
```

[...]

```
==1863== by 0x43F2FE: _ZN5boost9unit_test9framework3runEmb(framework.hpp:1629)  
==1863== by 0x463F10: _ZN5boost9unit_test14unit_test_mainEPFNSO_10test_suiteEiPPcEiS4_(unit_test_main.hpp:247)  
==1863== by 0x4648CB: main(unit_test_main.hpp:303)
```

[...]

Running 12 test cases...

```
acts-core/Tests/Utilities/BFieldMapUtilsTests.cpp(105): error: in "bfield_creation": difference{-nan} between  
value2_rz.perp(){-nan} and bfield2_rz.perp(){} exceeds 1e-09%
```